

Ridge_Regression.R

vijaykalmath

2022-01-04

```
# Ridge Regression and Lasso Regression

# GLMNET package is used for Lasso and Ridge Regression

# set seed for reproducibility
set.seed(1)

library(ISLR) # For the data
library(glmnet)

## Loading required package: Matrix

## Loaded glmnet 4.1-3

# Omitting NA
Hitters = na.omit(Hitters)

# Model Matrix produces matrix with 19 predictors but also transforms any qualitative variables into dummy variables
x = model.matrix(Salary~.,data=Hitters)[,-1]

y = Hitters$Salary

# Creating Exponential space gridspace for lambda or regularization factor
grid= 10^seq(10,-2,length=100)

# Alpha = 0 -> Ridge Regression , Alpha = 1 -> Lasso Regression

ridge.mod = glmnet(x,y,alpha=0,lambda = grid)

# 100 Models with corresponding Coefficients
dim(coef(ridge.mod))

## [1] 20 100

# Define Train and test set
train <- sample(1:nrow(x), nrow(x) / 2)
test <- (-train)
```

```
ridge.mod <- glmnet(x[train, ], y[train], alpha = 0,
                  lambda = grid, thresh = 1e-12)
```

```
ridge.pred <- predict(ridge.mod, s = 4, newx = x[test, ])
```

```
mean((ridge.pred - y[test])^2)
```

```
## [1] 142199.2
```

```
# With a single dimensional predictor , intercept only , we get the MSE of the constant mean predictor
mean((mean(y[train]) - y[test])^2)
```

```
## [1] 224669.9
```

```
# By making the lambda value very high , we can force the all the coeff to become closed to 0.
```

```
ridge.pred <- predict(ridge.mod, s = 1e10, newx = x[test, ])
```

```
mean((ridge.pred - y[test])^2)
```

```
## [1] 224669.8
```

```
ridge.pred <- predict(ridge.mod, s = 0, newx = x[test, ],
                    exact = T, x = x[train, ], y = y[train])
```

```
mean((ridge.pred - y[test])^2)
```

```
## [1] 168588.6
```

```
lm(y ~ x, subset = train)
```

```
##
```

```
## Call:
```

```
## lm(formula = y ~ x, subset = train)
```

```
##
```


```
## Coefficients:
```

## (Intercept)	xAtBat	xHits	xHmRun	xRuns	xRBI
## 274.0145	-0.3521	-1.6377	5.8145	1.5424	1.1243
## xWalks	xYears	xCAtBat	xCHits	xCHmRun	xCRuns
## 3.7287	-16.3773	-0.6412	3.1632	3.4008	-0.9739
## xCRBI	xWalks	xLeagueN	xDivisionW	xPutOuts	xAssists
## -0.6005	0.3379	119.1486	-144.0831	0.1976	0.6804
## xErrors	xNewLeagueN				
## -4.7128	-71.0951				

```
predict(ridge.mod, s = 0, exact = T, type = "coefficients",
       x = x[train, ], y = y[train])[1:20, ]
```

```
## (Intercept)      AtBat      Hits      HmRun      Runs      RBI
## 274.0200994    -0.3521900   -1.6371383   5.8146692   1.5423361   1.1241837
##      Walks      Years      CAtBat      CHits      CHmRun      CRuns
##   3.7288406  -16.3795195   -0.6411235   3.1629444   3.4005281  -0.9739405
##      CRBI      CWalks      LeagueN      DivisionW      PutOuts      Assists
##  -0.6003976    0.3378422  119.1434637 -144.0853061   0.1976300   0.6804200
##      Errors      NewLeagueN
##  -4.7127879  -71.0898914
```

```
cv.out <- cv.glmnet(x[train, ], y[train], alpha = 0)
plot(cv.out)
```



```
bestlam <- cv.out$lambda.min
```

```
bestlam
```

```
## [1] 431.0623
```

```
ridge.pred <- predict(ridge.mod, s = bestlam,
                      newx = x[test, ])
```

```
mean((ridge.pred - y[test])^2)
```

```
## [1] 138796
```

```
out <- glmnet(x, y, alpha = 0)
predict(out, type = "coefficients", s = bestlam)[1:20, ]
```

##	(Intercept)	AtBat	Hits	HmRun	Runs	RBI
##	24.88203033	0.09571573	0.77893014	0.85488974	1.02412478	0.87575813
##	Walks	Years	CAtBat	CHits	CHmRun	CRuns
##	1.51400163	1.93680144	0.01128321	0.05331300	0.38052103	0.10651815
##	CRBI	CWalks	LeagueN	DivisionW	PutOuts	Assists
##	0.11215778	0.06298811	18.94995835	-70.27086710	0.14893846	0.02329106
##	Errors	NewLeagueN				
##	-1.09701573	9.45875165				